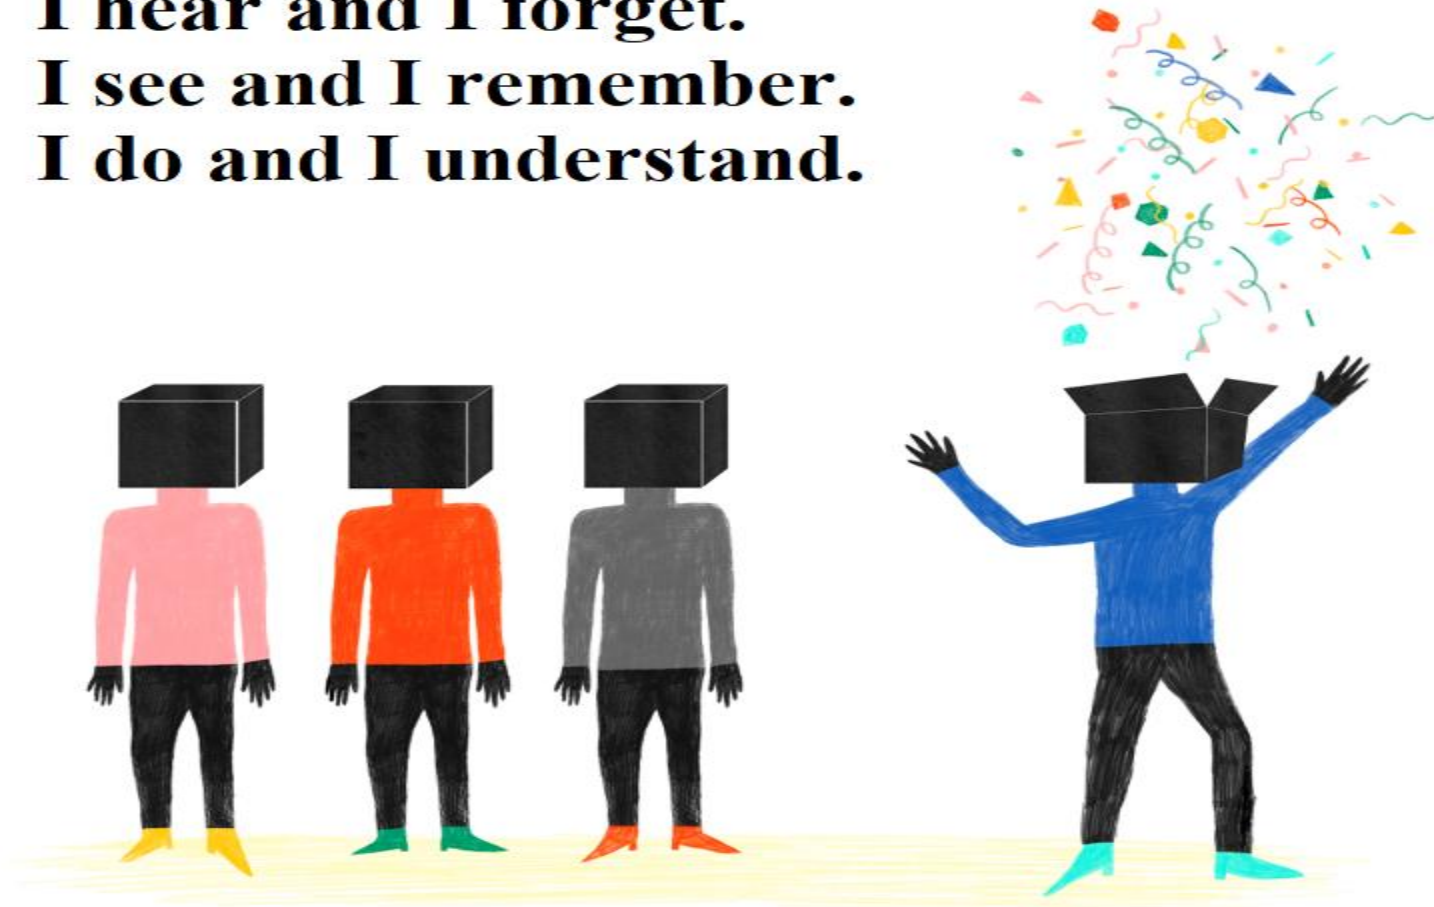


**I hear and I forget.
I see and I remember.
I do and I understand.**




Chandan Verma


Corporate Trainer(Machine Learning,AI,Cloud Computing,IOT)

www.facebook.com/verma.chandan.070

Curve fitting

Curve fitting, essentially, is similar to the game of connecting the dots where you try to complete a picture





0	2.0	4	8.5	196
1	2.4	4	9.6	221
2	1.5	4	5.9	136
3	3.5	6	11.1	255
4	3.5	6	10.6	244
5	3.5	6	10.0	230
6	3.5	6	10.1	232
7	3.7	6	11.1	255
8	3.7	6	11.6	267
9	2.4	4	9.2	?

Regression is the process of predicting a Continuous Value


Linear Regression

Linear regression analysis is used to predict the value of a variable based on the value of another variable. The variable you want to predict is called the dependent variable. The variable you are using to predict the other variable's value is called the independent variable.


Regression is a statistical way to establish a relationship between a dependent variable and a set of independent variable(s).

Regression

Year	Sales (Million Euro)	Advertising (Million Euro)
1	651	23
2	762	26
3	856	30
4	1,063	34
5	1,190	43
6	1,298	48
7	1,421	52
8	1,440	57
9	1,518	58



Linear Regression

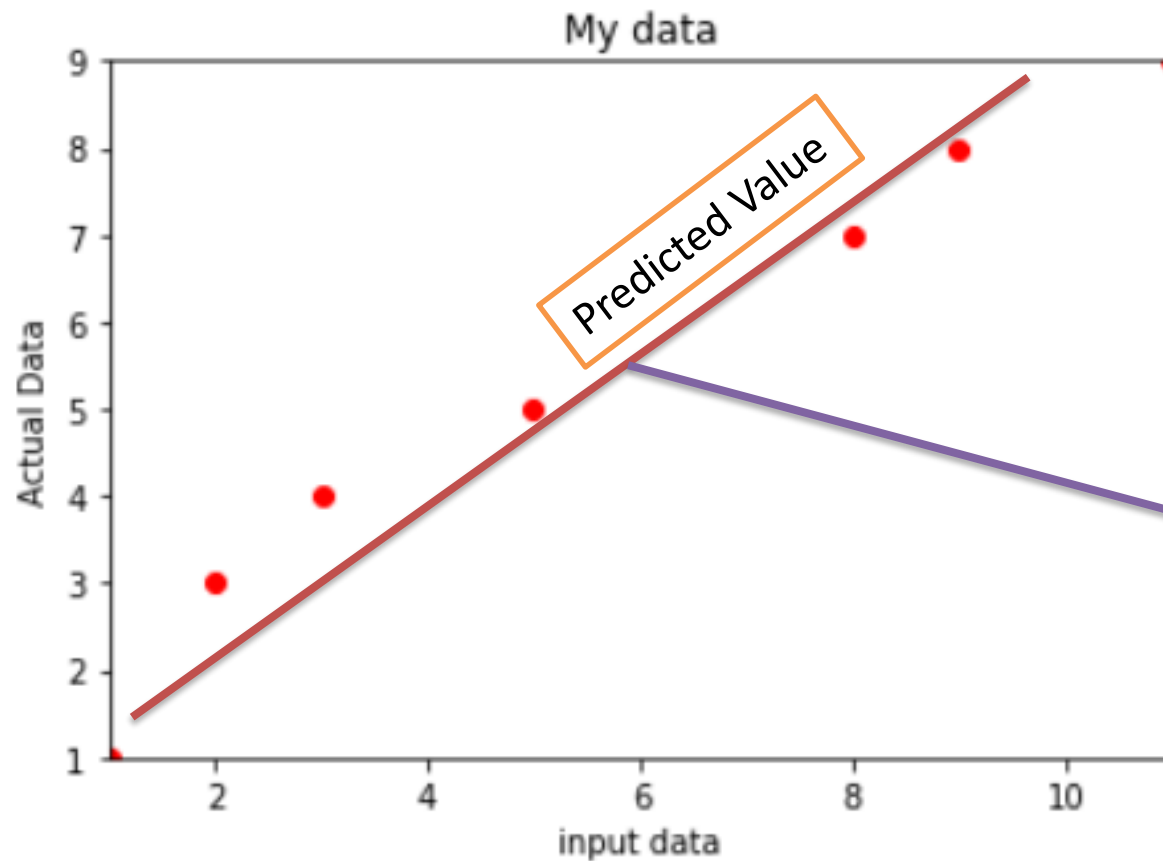


Linear Regression

Linear regression is a statistical approach for modeling relationship between a **dependent(response)** variable with a given set of **independent(features)** variables.

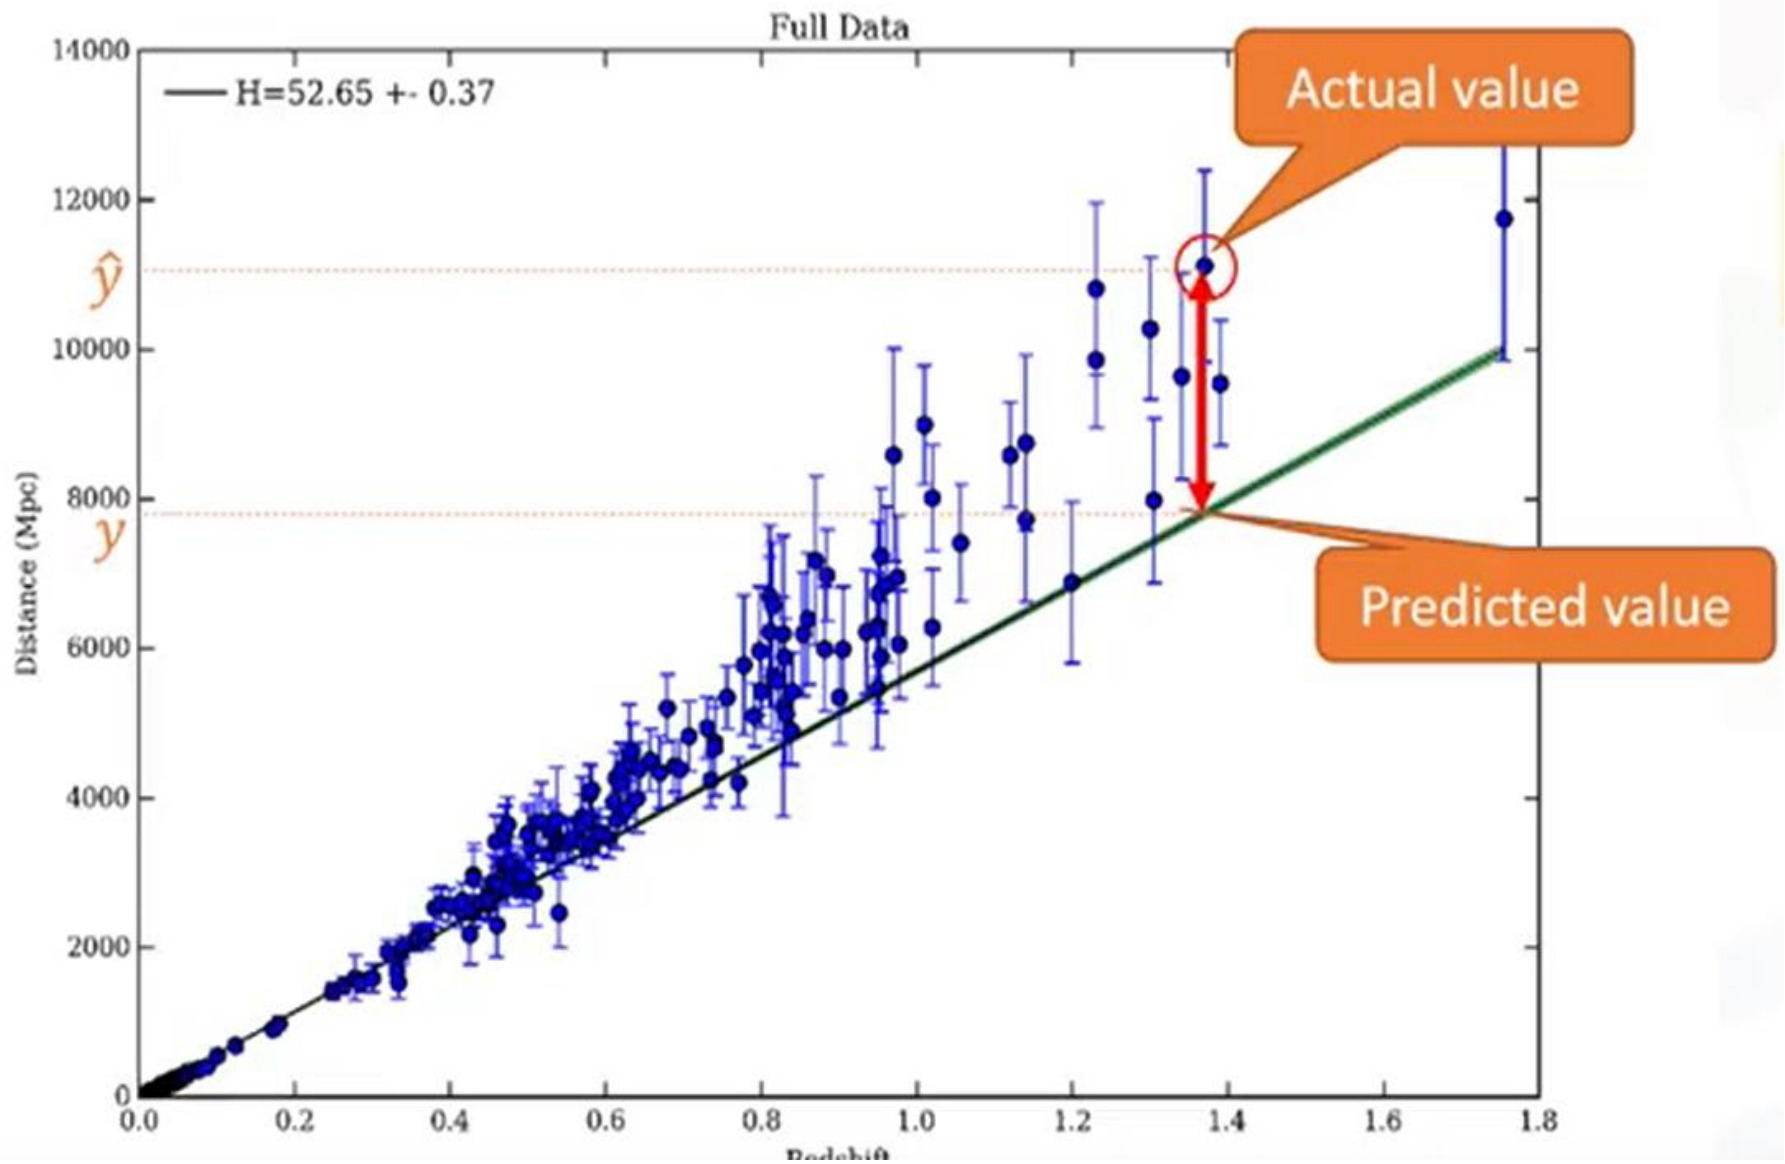
special case of fitting a straight line to some given data


In statistics, linear regression is an approach for modeling the relationship between a scalar dependent variable y and one or more explanatory variables (or independent variables) denoted X .





**Best fit Line
Minimize Loss**

$$\text{Loss} = (\text{Actual} - \text{Predicted})$$






It is a method to predict the target variable by finding a best fit line between the independent and dependent variable. The best fit is the line with minimum error from all the points.



The technique we will use to find the *best fitting* line will be called the *method of least squares*.



$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2.$$

Best fit line



The diagram illustrates the Mean Absolute Error (MAE) formula with color-coded annotations:


- Divide by the total number of data points:** A blue line points to the $\frac{1}{n}$ term, which is enclosed in a blue box.
- Sum of:** A black line points to the summation symbol Σ .
- Actual output value:** A green line points to the y term, which is enclosed in a green box.
- Predicted output value:** A yellow line points to the \hat{y} term, which is enclosed in an orange box.
- The absolute value of the residual:** A black bracket underneath the $|y - \hat{y}|$ term is labeled with this text.


$$MAE = \frac{1}{n} \sum |y - \hat{y}|$$

There are simple linear regression calculators that use a “least squares” method to discover the best-fit line for a set of paired data. You then estimate the value of X (dependent variable) from Y (independent variable)


Best fit line







x	y
1	1
2	3
4	3
3	2
5	5



$$y = B0 + B1 * x$$

Simple Linear Regression

1. Mean Function.
2. Variance Function.
3. Covariance Function.
4. Functions to calculate the w_0 and w_1 values

$$w_1 = \frac{\text{covariance}(x,y)}{\text{variance}(x)}$$

$$w_0 = \text{mean}(y) - (w_1 * \text{mean}(x))$$

Simple Linear Regression


Having one independent variable to predict the dependent variable.

$$\sigma^2 = \frac{\sum_{i=1}^n (x_i - \text{mean}(x))^2}{n-1}$$

$$\text{COV}_{x,y} = \frac{\sum_{i=1}^N (x_i - \text{mean}(x))(y_i - \text{mean}(y))}{N-1}$$


$$\text{mean}(x) = \frac{(x_1) + (x_2) + (x_3) \dots + (x_n)}{n}$$

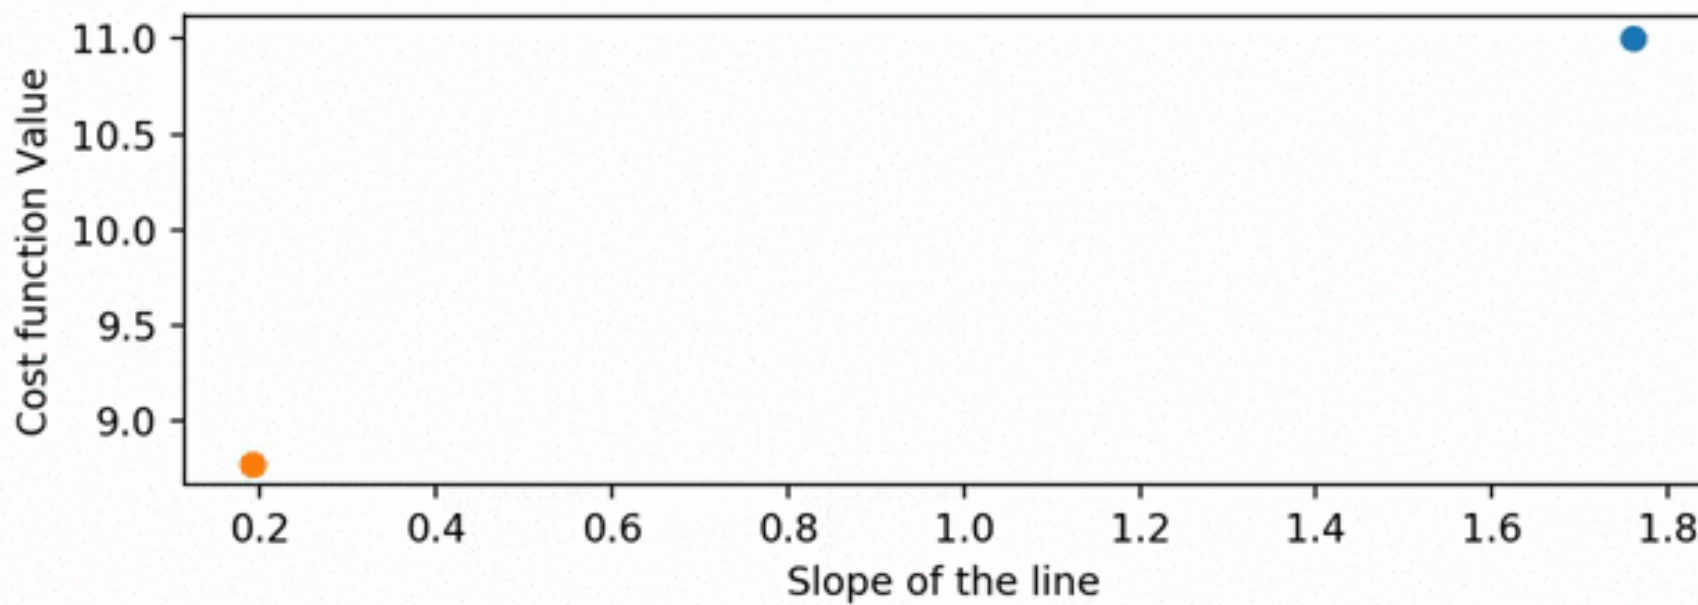
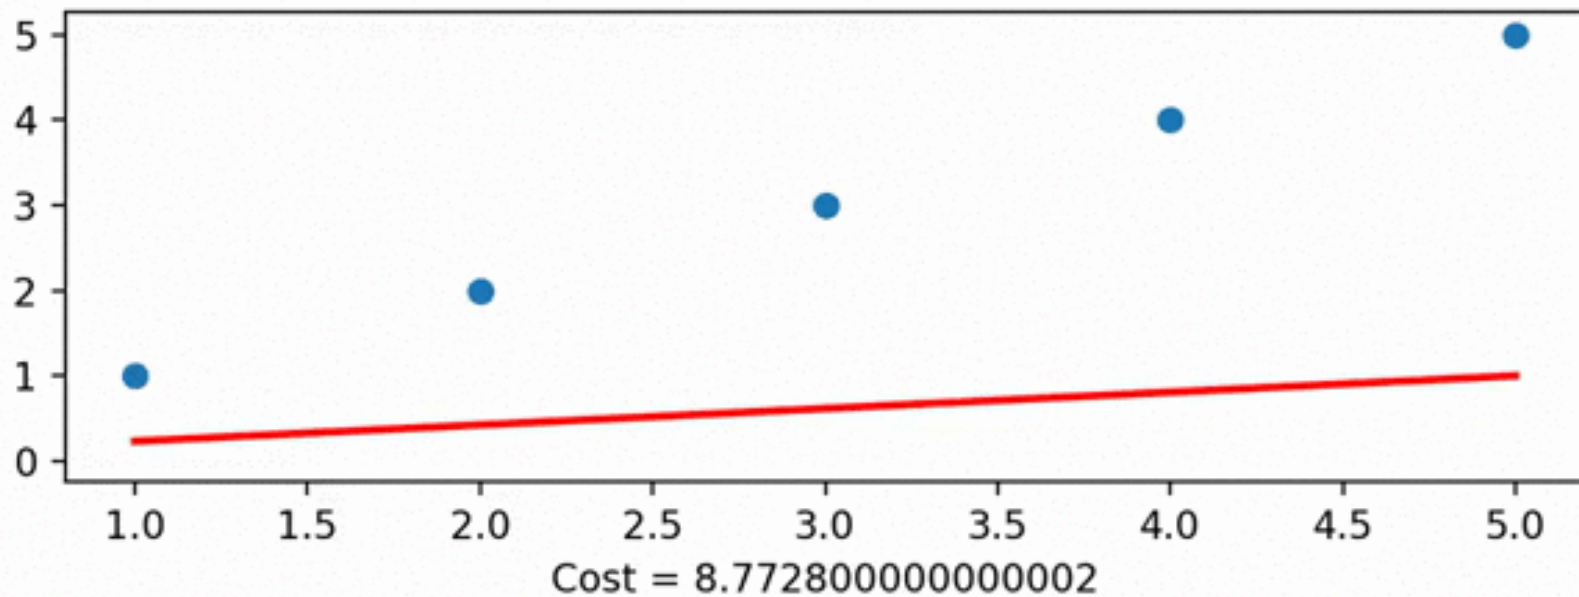
$y(x) = ax + b$ where a and b are unknown real values





$$y(x) = .8 x + .4$$

$$y(x) = 1.3 x + 2.5$$









Cost Function

cost function is often the squared of the difference between actual and predicted outcome, because the difference can sometimes be negative; this is also known as *min least-squared*

$$\text{minimize } \frac{1}{n} \sum_{i=1}^n (\text{pred}_i - y_i)^2$$

Error of the Model

$$MAE = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$RMSE = \sqrt{\frac{1}{n} \sum_{j=1}^n (y_j - \hat{y}_j)^2}$$

$$RAE = \frac{\sum_{j=1}^n |y_j - \hat{y}_j|}{\sum_{j=1}^n |y_j - \bar{y}|}$$

$$RSE = \frac{\sum_{j=1}^n (y_j - \hat{y}_j)^2}{\sum_{j=1}^n (y_j - \bar{y})^2}$$

$$R^2 = 1 - RSE$$

Problem

X (sleep, study)	y (test score)
(3,5)	75
(5,1)	82
(10,2)	93
(8,3)	?

we want to predict our test score based on how many hours we sleep and how many hours we study the night before


Finding Minimum Value of a Function

$$f(x) : x^4 - 10x^2 + 9$$

$$\frac{d}{dx}f(x) = 4x^3 - 20x = 0$$

$$\text{Minima } f(x) : x = \pm\sqrt{5} = \pm 2.24$$

Iterative Calculation



Gradient Descent

Gradient descent is an optimization algorithm used to find the values of parameters (coefficients) of a function (f) that minimizes a cost function (cost).

$$\text{Error}_{(m,b)} = \frac{1}{N} \sum_{i=1}^N (y_i - (mx_i + b))^2$$

$$\frac{\partial}{\partial m} = \frac{2}{N} \sum_{i=1}^N -x_i (y_i - (mx_i + b))$$

$$\frac{\partial}{\partial b} = \frac{2}{N} \sum_{i=1}^N -(y_i - (mx_i + b))$$

Gradient Descent

$$\text{Loss Function } (LF) = \frac{1}{N} \sum (y - (mx + c))^2$$


$$\frac{\partial}{\partial m} LF = \frac{2}{N} \sum (y - (mx + c)) \times x$$

$$\frac{\partial}{\partial c} LF = \frac{2}{N} \sum (y - (mx + c))$$

when our model coefficients are trained. During training, our initial weights are updated according to a gradient update rule using a learning rate and a gradient

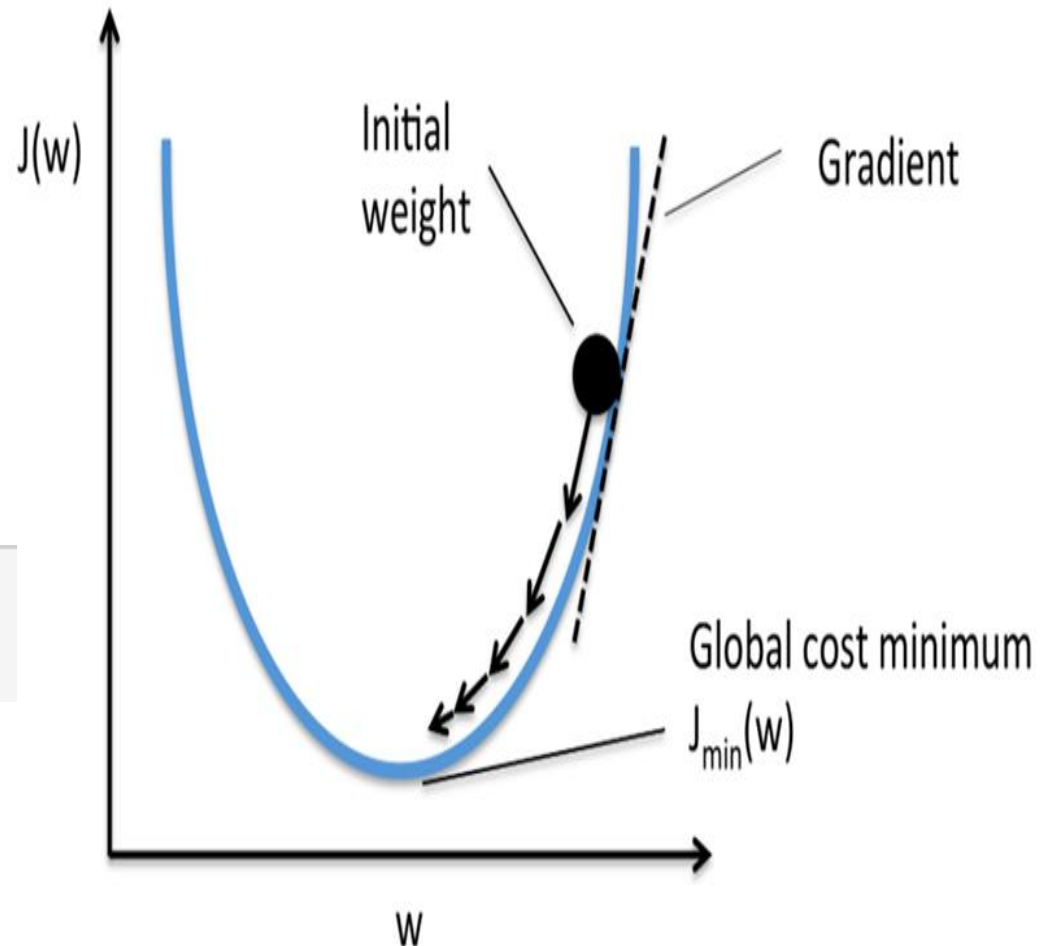

Iterative Calculation


Repeat until converge




$$m_{n+1} = m_n - \alpha \frac{\partial}{\partial m_n} LF(m_n)$$

$$c_{n+1} = c_n - \alpha \frac{\partial}{\partial c_n} LF(c_n)$$








repeat until convergence {

$$\theta_0 := \theta_0 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})$$


$$\theta_1 := \theta_1 - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

}


Gradient descent of x^2 , learning rate=0.2





Loss vs. Weight



Gradient Descent







X_INPUT	Y_OUTPUT
1	6
2	5
3	7
4	10

$$y = \beta_1 + \beta_2 x \quad \beta_1 + 1\beta_2 = 6$$

$$\beta_1 + 2\beta_2 = 5$$

$$\beta_1 + 3\beta_2 = 7$$

$$\beta_1 + 4\beta_2 = 10$$

$$\begin{aligned}
 S(\beta_1, \beta_2) &= [6 - (\beta_1 + 1\beta_2)]^2 + [5 - (\beta_1 + 2\beta_2)]^2 \\
 &\quad + [7 - (\beta_1 + 3\beta_2)]^2 + [10 - (\beta_1 + 4\beta_2)]^2 \\
 &= 4\beta_1^2 + 30\beta_2^2 + 20\beta_1\beta_2 - 56\beta_1 - 154\beta_2 + 210.
 \end{aligned}$$

The minimum is determined by calculating the partial derivatives of $S(\beta_1, \beta_2)$ with respect to β_1 and β_2 and setting them to zero

$$\frac{\partial S}{\partial \beta_1} = 0 = 8\beta_1 + 20\beta_2 - 56$$

$$\beta_1 = 3.5$$

$$\frac{\partial S}{\partial \beta_2} = 0 = 20\beta_1 + 60\beta_2 - 154.$$

$$\beta_2 = 1.4$$

equation $y = 3.5 + 1.4x$ of the line of best fit

The diagram illustrates the components of the linear regression equation $Y_i = \beta_0 + \beta_1 X_i + \epsilon_i$. The equation is written in large, bold, black font. Above the equation, four labels with arrows point to specific parts: 'Dependent Variable' points to Y_i , 'Population Y intercept' points to β_0 , 'Population Slope Coefficient' points to β_1 , and 'Independent Variable' points to X_i . Below the equation, two blue curly braces are used: one under $\beta_0 + \beta_1 X_i$ labeled 'Linear component', and another under ϵ_i labeled 'Random Error component'. To the right of the equation, the label 'Random Error term' has an arrow pointing to ϵ_i .

Regression-model

$$\boxed{\hat{y}} = \beta_0 + \beta_1 \boxed{x} + \boxed{\epsilon}$$

Predicted output

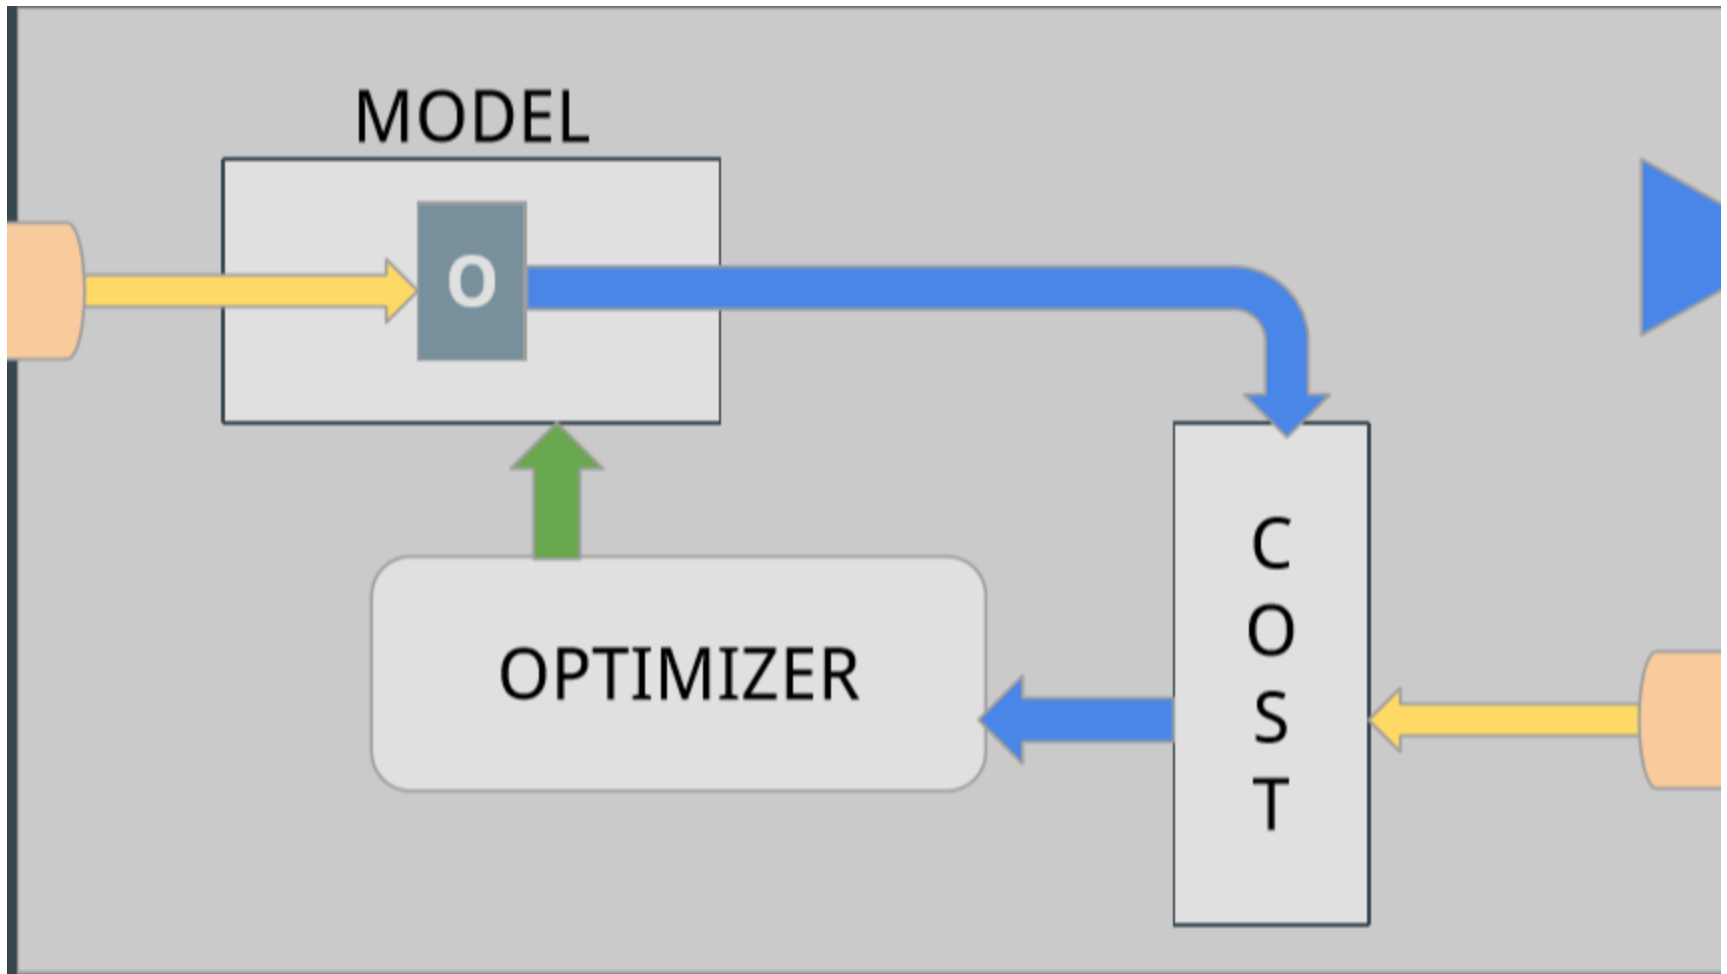
Coefficients

Input

Error

$$\boxed{\hat{y}} = \beta_0 + \beta_1 \boxed{x_1} + \dots + \beta_p \boxed{x_p} + \boxed{\epsilon}$$

Regression



Cost (loss) function

$$\mathbf{x} \cdot \mathbf{W} + \mathbf{b} = \mathbf{y}(\text{target})$$

prediction

$$\mathbf{y} - (\mathbf{x} \cdot \mathbf{W} + \mathbf{b})$$

$$[\mathbf{y} - (\mathbf{x} \cdot \mathbf{W} + \mathbf{b})]^2$$

$$\Sigma [\mathbf{y}_i - (\mathbf{x}_i \cdot \mathbf{W} + \mathbf{b})]^2$$

$$\text{cost} = \text{tf.reduce_sum}(\text{tf.square}(\text{output} - \mathbf{y}))$$

Assumptions of linear regression

1. There must be a linear relationship between the dependent and independent variables.
2. Error terms are normally distributed with mean 0.
3. No multicollinearity - When the independent variables in my model are highly linearly related then such a situation is called multicollinearity.
4. No outliers are present in the data.

sklearn.linear_model.LinearRegression

- `fit_intercept=True`,
- `normalize=False`,
- `copy_X=True`,
- `n_jobs=None`)

variance, r² score, and mean square error

It is important to understand these metrics to determine whether regression models are accurate or misleading.

variance—in terms of linear regression, **variance** is a measure of how far observed values differ from the average of predicted values, i.e., their difference from the **predicted value mean**.

r² score—varies between 0 and 100%. It is closely related to the **MSE**, but not the same

R square

It suggests the proportion of variation in Y which can be explained with the independent variables. Mathematically, it is the ratio of predicted values and observed values

R-Squared is a statistical measure of fit that indicates how much variation of a dependent variable is explained by the independent variable(s) in a regression model.

$$R^2 = 1 - \frac{\text{Explained Variation}}{\text{Total Variation}}$$

R-squared can take value between 0 and 1 where values closer to 0 represent a poor fit while values closer to 1 represent a perfect fit

$$R^2 = \frac{SSR}{SST} = \frac{\sum (\hat{y}_i - \bar{y})^2}{\sum (y_i - \bar{y})^2}$$

$$R - Square = 1 - \frac{\sum (Y_{actual} - Y_{predicted})^2}{\sum (Y_{actual} - Y_{mean})^2}$$

$$R^2 = \frac{\sum y_i^2 - \sum e_i^2}{\sum y_i^2} = 1 - \frac{\sum e_i^2}{\sum y_i^2}$$

Adjusted R-square

The only drawback of R^2 is that if new predictors (X) are added to our model, R^2 only increases or remains constant but it never decreases. We can not judge that by increasing complexity of our model, are we making it more accurate?

$$R^2 \text{ adjusted} = 1 - \frac{(1 - R^2)(N - 1)}{N - p - 1}$$

gradientdescentoptimizer

Gradient Descent is a learning algorithm that attempts to minimise some error

performing gradient descent to minimize the cost function, to obtain the 'good' values for W, b.

```
optimizer =  
tf.train.GradientDescentOptimizer(learning_rate=0.00001)
```

```
optimizer = optimizer.minimize(cost)
```

The learning rate is how quickly a network **abandons** old beliefs for new ones

Linear Regression Exercise

- Creating training data
- Placeholders
- Modeling
- Training

Linear Model (in TF notation):

$$y = \text{tf.matmul}(x, W) + b$$

The goal in linear regression is to find W , b , such that given any feature value (x), we can find the prediction (y) by substituting W , x , b values into the model.

need to define a cost function, which is a measure of the *difference* between the **prediction** (y) for given a feature value (x), and the **actual outcome** (y_{actual}) for that same feature value (x).

Model linear regression $y = Wx + b$

$$[y - (x \cdot W + b)]^2$$

```
tf.train.GradientDescentOptimizer(LR).minimize(cost)
```

```
cost = tf.reduce_mean(tf.pow((y_blue - yred), 2))
```

```
yred = tf.matmul(x, W) + b
```

```
y_blue = tf.placeholder(tf.float32, [None, 1])
```

```
x = tf.placeholder(tf.float32, [None, 1])
```

```
import numpy as np
import tensorflow as tf
x = tf.placeholder(tf.float32, [None, 1])
W = tf.Variable(tf.zeros([1,1]))
b = tf.Variable(tf.zeros([1]))
product = tf.matmul(x,W)
y = product + b
y_ = tf.placeholder(tf.float32, [None, 1])
cost = tf.reduce_mean(tf.square(y_-y))
train_step =
tf.train.GradientDescentOptimizer(0.0000001).minimize(cost)
```



```
sess = tf.Session()
init = tf.initialize_all_variables()
sess.run(init)
steps = 1000
for i in range(steps):
    xs = np.array([[i]])
    ys = np.array([[2*i]])
    feed = { x: xs, y_: ys }
    sess.run(train_step, feed_dict=feed)
    print("After %d iteration:" % i)
    print("W: %f" % sess.run(W))
    print("b: %f" % sess.run(b))
    print("cost: %f" % sess.run(cost, feed_dict=feed))
```